


Computer Architecture




Building a Modern Computer From First Principles



www.nand2tetris.org

Babbage's Analytical Engine (1835)

"We may say most aptly that the Analytical Engine weaves algebraic patterns just as the Jacquard-loom weaves flowers and leaves"
(Ada Lovelace)




Charles Babbage (1791-1871)





Some early computers and computer scientists




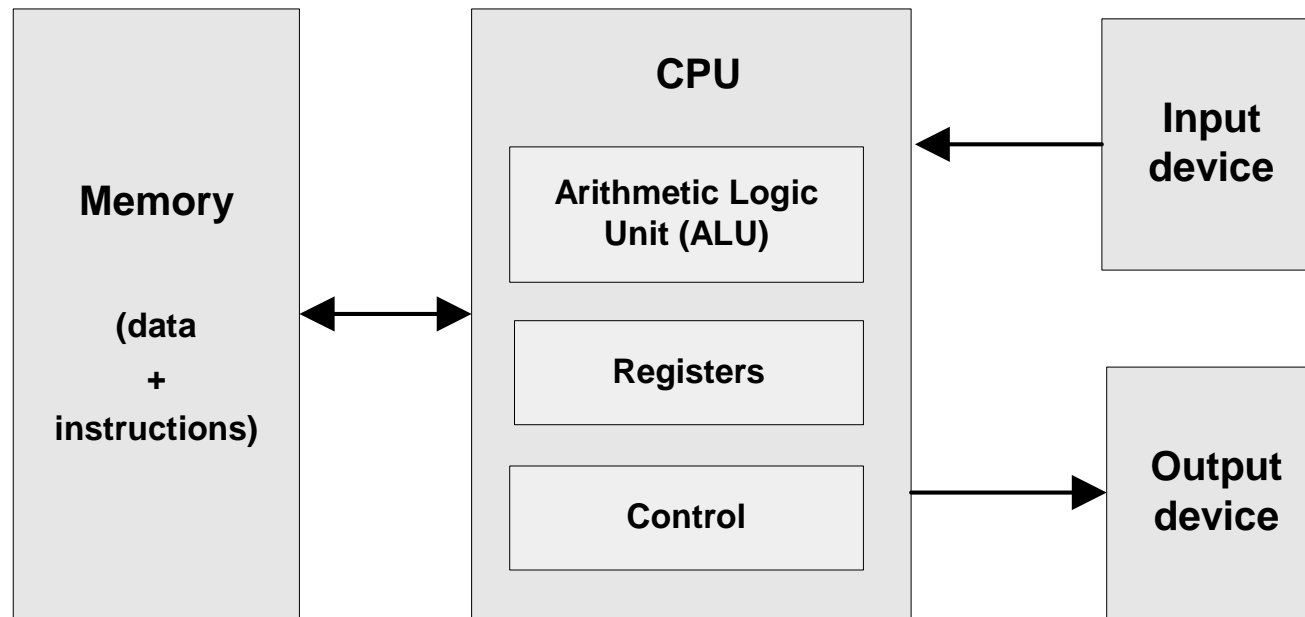
Blaise Pascal
1623-1662



Gottfried Leibniz
1646-1716



Von Neumann machine (circa 1940)




John Von Neumann (and others) ... made it possible

Stored
program
concept!



Andy Grove (and others) ... made it small and fast.

Processing logic: fetch-execute cycle



Executing the *current instruction* involves one or more of the following micro-tasks:

- ❑ Have the ALU compute some function $out = f(\text{register values})$
- ❑ Write the ALU output to selected registers
- ❑ As a side-effect of this computation, figure out which instruction to fetch and execute next.

The Hack chip-set and hardware platform

Elementary logic gates

- Nand
- Not
- And
- Or
- Xor
- Mux
- Dmux
- Not16
- And16
- Or16
- Mux16
- Or8Way
- Mux4Way16
- Mux8Way16
- DMux4Way
- DMux8Way

done

Combinational chips

- HalfAdder
- FullAdder
- Add16
- Inc16
- ALU

done

Sequential chips

- DFF
- Bit
- Register
- RAM8
- RAM64
- RAM512
- RAM4K
- RAM16K
- PC

done

Computer Architecture

- Memory
- CPU
- Computer

this lecture


The Hack computer

- A 16-bit Von Neumann platform
- The *instruction memory* and the *data memory* are physically separate
- Screen: 512 rows by 256 columns, black and white
- Keyboard: standard
- Designed to execute programs written in the Hack machine language
- Can be easily built from the chip-set that we built so far in the course

Main parts of the Hack computer:

- ❑ Instruction memory (ROM)
- ❑ Memory (RAM):
 - Data memory
 - Screen (memory map)
 - Keyboard (memory map)
- ❑ CPU
- ❑ Computer (the logic that holds everything together).

Lecture / construction plan



- Instruction memory

- Memory:

- Data memory


- Screen

- Keyboard

- CPU

- Computer

Instruction memory



Function:

- The ROM is pre-loaded with a program written in the Hack machine language
- The ROM chip always emits a 16-bit number:

`out = ROM32K[address]`


- This number is interpreted as the *current instruction*.

Data memory

Low-level (hardware) read/write logic:

To read $\text{RAM}[k]$: set address to k ,
probe out

To write $\text{RAM}[k]=x$: set address to k ,
set in to x ,
set load to 1,
run the clock



High-level (OS) read/write logic:

To read $\text{RAM}[k]$: use the OS command $\text{out} = \text{peek}(k)$

To write $\text{RAM}[k]=x$: use the OS command $\text{poke}(k, x)$


peek and poke are OS commands whose implementation should effect the same behavior as the low-level commands

More about peek and poke this later in the course, when we'll write the OS.

Lecture / construction plan

- ✓ ■ Instruction memory
 - Memory:
 - ✓ □ Data memory
 - □ Screen
 - Keyboard
 - CPU
 - Computer

Screen



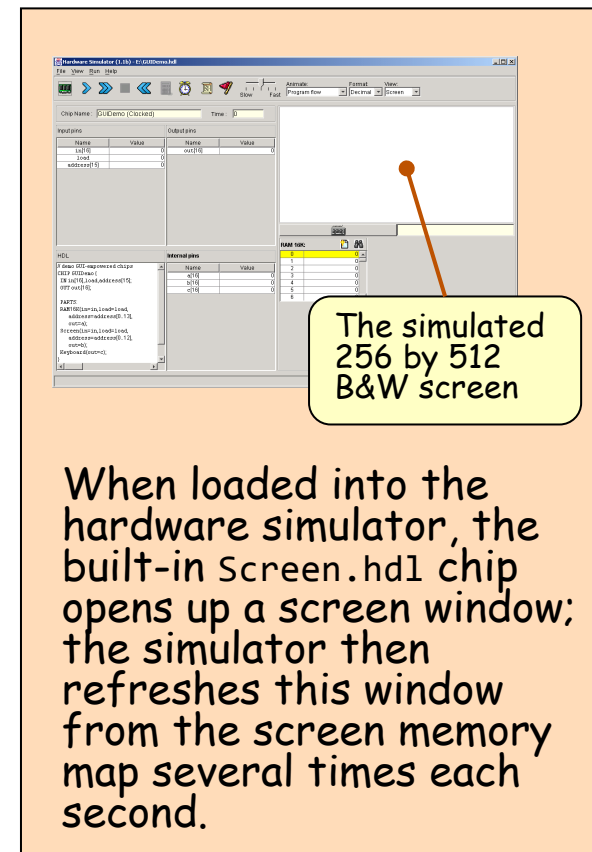
The Screen chip has a basic RAM chip functionality:

- read logic: $out = Screen[address]$
- write logic: if load then $Screen[address] = in$

Side effect:

Continuously refreshes a 256 by 512 black-and-white screen device

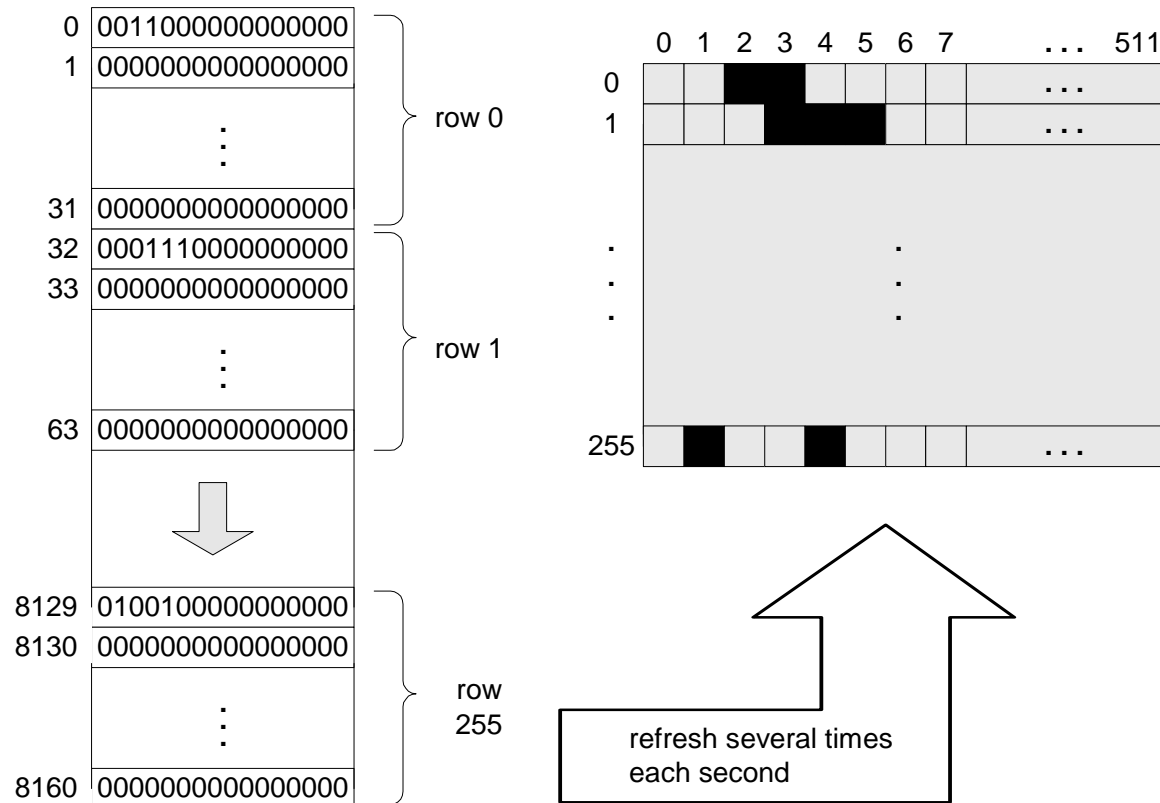

Simulated screen:



When loaded into the hardware simulator, the built-in Screen.hdl chip opens up a screen window; the simulator then refreshes this window from the screen memory map several times each second.

Screen memory map


In the Hack platform,
the screen is
implemented as an 8K
16-bit RAM chip.



How to set the (row,col) pixel of the screen to black or to white:

- ❑ Low-level (machine language): Set the $col\%16$ bit of the word found at $Screen[row*32+col/16]$ to 1 or to 0 ($col/16$ is integer division)
- ❑ High-level: Use the OS command `drawPixel(row,col)` (effects the same operation, discussed later in the course, when we'll write the OS).

Keyboard



Keyboard chip: a single 16-bit register

Input: scan-code (16-bit value) of the currently pressed key, or 0 if no key is pressed

Output: same

<u>Special keys:</u>				
Key pressed	Keyboard output		Key pressed	Keyboard output
newline	128		end	135
backspace	129		page up	136
left arrow	130		page down	137
up arrow	131		insert	138
right arrow	132		delete	139
down arrow	133		esc	140
home	134		f1-f12	141-152

How to read the keyboard:

- ❑ Low-level (hardware): probe the contents of the Keyboard chip
- ❑ High-level: use the OS command `keyPressed()`
(effects the same operation, discussed later in the course, when we'll write the OS).

Simulated keyboard:

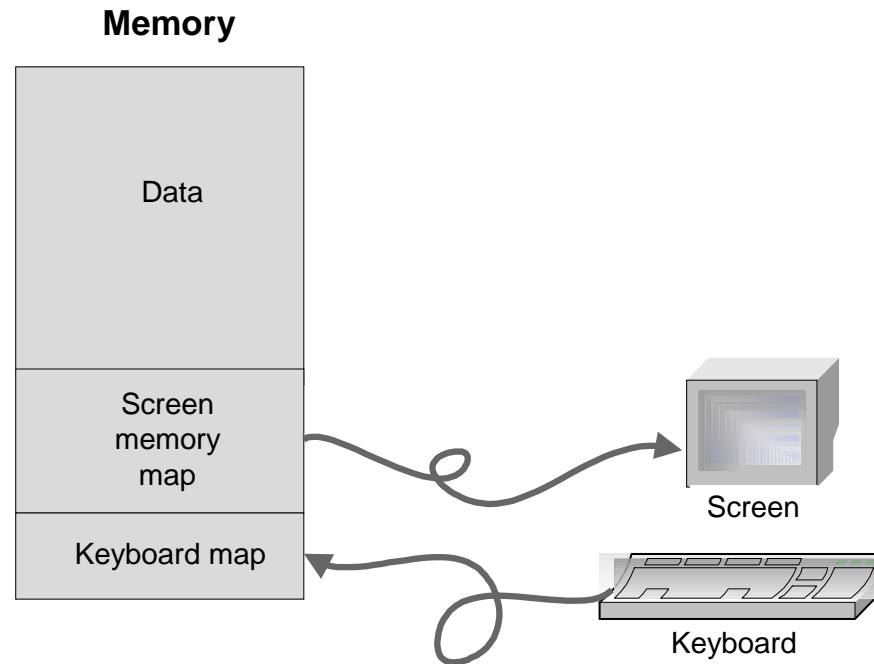
The screenshot shows a logic simulator interface with a 'Keyboard' chip selected. A callout box points to a button labeled 'The simulated keyboard enabler button'. Below the screenshot, text explains the implementation.

The keyboard is implemented as a built-in `keyboard.hdl` chip. When this java chip is loaded into the simulator, it connects to the regular keyboard and pipes the scan-code of the currently pressed key to the keyboard memory map.

Lecture / construction plan

- ✓ ■ Instruction memory
- ➡ ■ Memory:
 - ✓ □ Data memory
 - ✓ □ Screen
 - ✓ □ Keyboard
- CPU
- Computer

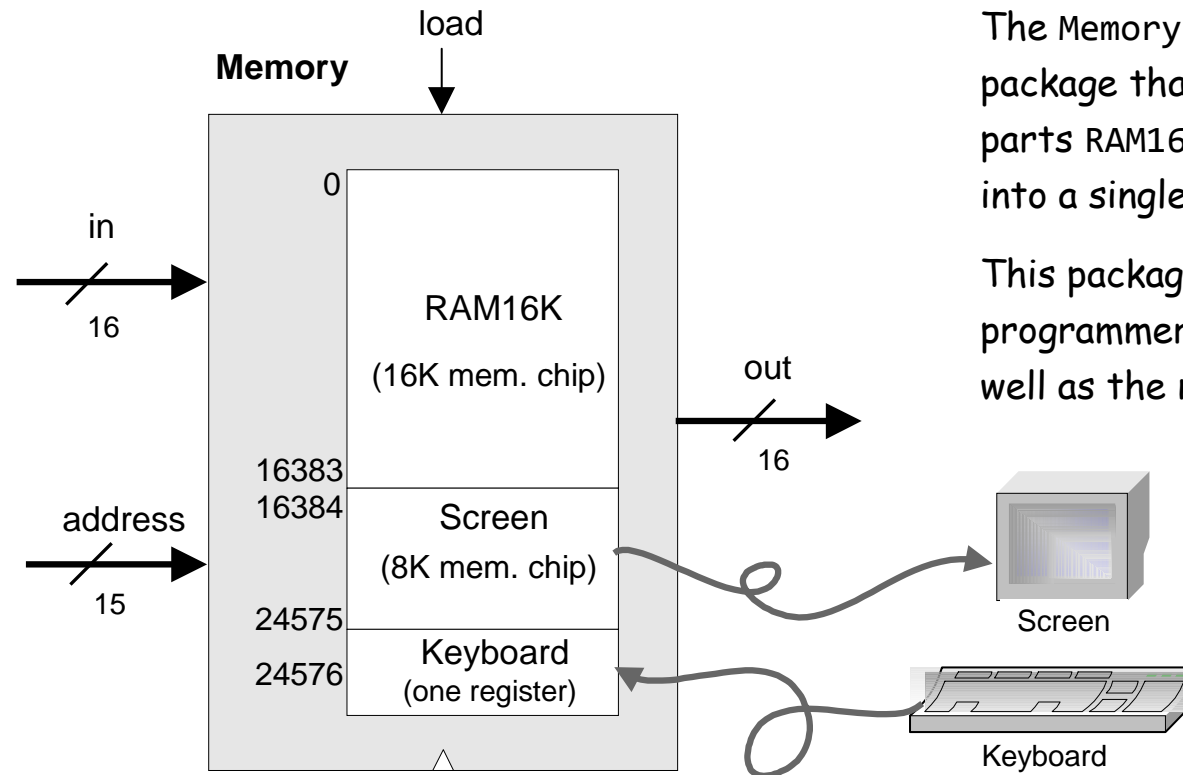
Memory: conceptual / programmer's view



Using the memory:

- ❑ To record or recall values (e.g. variables, objects, arrays), use the first 16K words of the memory
- ❑ To write to the screen (or read the screen), use the next 8K words of the memory
- ❑ To read which key is currently pressed, use the next word of the memory.

Memory: physical implementation



The Memory chip is essentially a package that integrates the three chip-parts RAM16K, Screen, and Keyboard into a single, contiguous address space.

This packaging effects the programmer's view of the memory, as well as the necessary I/O side-effects.

Access logic:

- ❑ Access to any address from 0 to 16,383 results in accessing the RAM16K chip-part
- ❑ Access to any address from 16,384 to 24,575 results in accessing the Screen chip-part
- ❑ Access to address 24,576 results in accessing the keyboard chip-part
- ❑ Access to any other address is invalid.


Lecture / construction plan

- ✓ ■ Instruction memory
- ✓ ■ Memory:
 - ✓ ■ Data memory
 - ✓ ■ Screen
 - ✓ ■ Keyboard
- ➡ ■ CPU
- Computer


A pledge to patience ...

"At times ... the fragments that I lay out for your inspection may seem not to fit well together, as if they were stray pieces from separate puzzles. In such cases, I would counsel patience. There are moments when a large enough fragment can become a low wall, a second fragment another wall to be raised at a right angle to the first. A few struts and beams later, and we may have made ourselves a rough foundation ... But it can consume the better part of a chapter to build such a foundation; and as we do so the fragment that we are examining may seem unconnected to the larger whole. Only when we step back can we see that we have been assembling something that can stand in the wind."

From: *Sailing the Wind Dark Sea* (Thomas Cahill)



CPU




CPU internal components (invisible in this chip diagram): ALU and 3 registers: A, D, PC

CPU execute logic:

The CPU executes the instruction according to the Hack language specification:

- ❑ The D and A values, if they appear in the instruction, are read from (or written to) the respective CPU-resident registers
- ❑ The M value, if there is one in the instruction's RHS, is read from inM
- ❑ If the instruction's LHS includes M, then the ALU output is placed in outM, the value of the CPU-resident A register is placed in addressM, and writeM is asserted.

CPU



CPU internal components (invisible in this chip diagram): ALU and 3 registers: A, D, PC

CPU fetch logic:

Recall that:

1. the instruction may include a jump directive (expressed as non-zero jump bits)
2. the ALU emits two control bits, indicating if the ALU output is zero or less than zero

If reset==0: the CPU uses this information (the jump bits and the ALU control bits) as follows:

If there should be a jump, the PC is set to the value of A; else, PC is set to PC+1

If reset==1: the PC is set to 0. (restarting the computer)

The C-instruction revisited

dest = comp; jump

comp

dest

jump

binary:

1 1 1 a

c1 c2 c3 c4

c5 c6 d1 d2

d3 j1 j2 j3

(when a=0) <i>comp</i>	c1	c2	c3	c4	c5	c6	(when a=1) <i>comp</i>	d1	d2	d3	Mnemonic	Destination (where to store the computed value)
0	1	0	1	0	1	0		0	0	0	null	The value is not stored anywhere
1	1	1	1	1	1	1		0	0	1	M	Memory[A] (memory register addressed by A)
-1	1	1	1	0	1	0		0	1	0	D	D register
D	0	0	1	1	0	0		0	1	1	MD	Memory[A] and D register
A	1	1	0	0	0	0	M	1	0	0	A	A register
!D	0	0	1	1	0	1		1	0	1	AM	A register and Memory[A]
!A	1	1	0	0	0	1	!M	1	1	0	AD	A register and D register
-D	0	0	1	1	1	1		1	1	1	AMD	A register, Memory[A], and D register
-A	1	1	0	0	1	1	-M					
D+1	0	1	1	1	1	1						
A+1	1	1	0	1	1	1	M+1					
D-1	0	0	1	1	1	0						
A-1	1	1	0	0	1	0	M-1					
D+A	0	0	0	0	1	0	D+M					
D-A	0	1	0	0	1	1	D-M					
A-D	0	0	0	1	1	1	M-D					
D&A	0	0	0	0	0	0	D&M					
D A	0	1	0	1	0	1	D M					

j1 (out < 0)	j2 (out = 0)	j3 (out > 0)	Mnemonic	Effect
0	0	0	null	No jump
0	0	1	JGT	If out > 0 jump
0	1	0	JEQ	If out = 0 jump
0	1	1	JGE	If out ≥ 0 jump
1	0	0	JLT	If out < 0 jump
1	0	1	JNE	If out ≠ 0 jump
1	1	0	JLE	If out ≤ 0 jump
1	1	1	JMP	Jump

CPU implementation

dest = comp; jump

comp

dest

jump

binary:

1

1

1

a

c1

c2

c3

c4

c5

c6

d1

d2

d3

j1

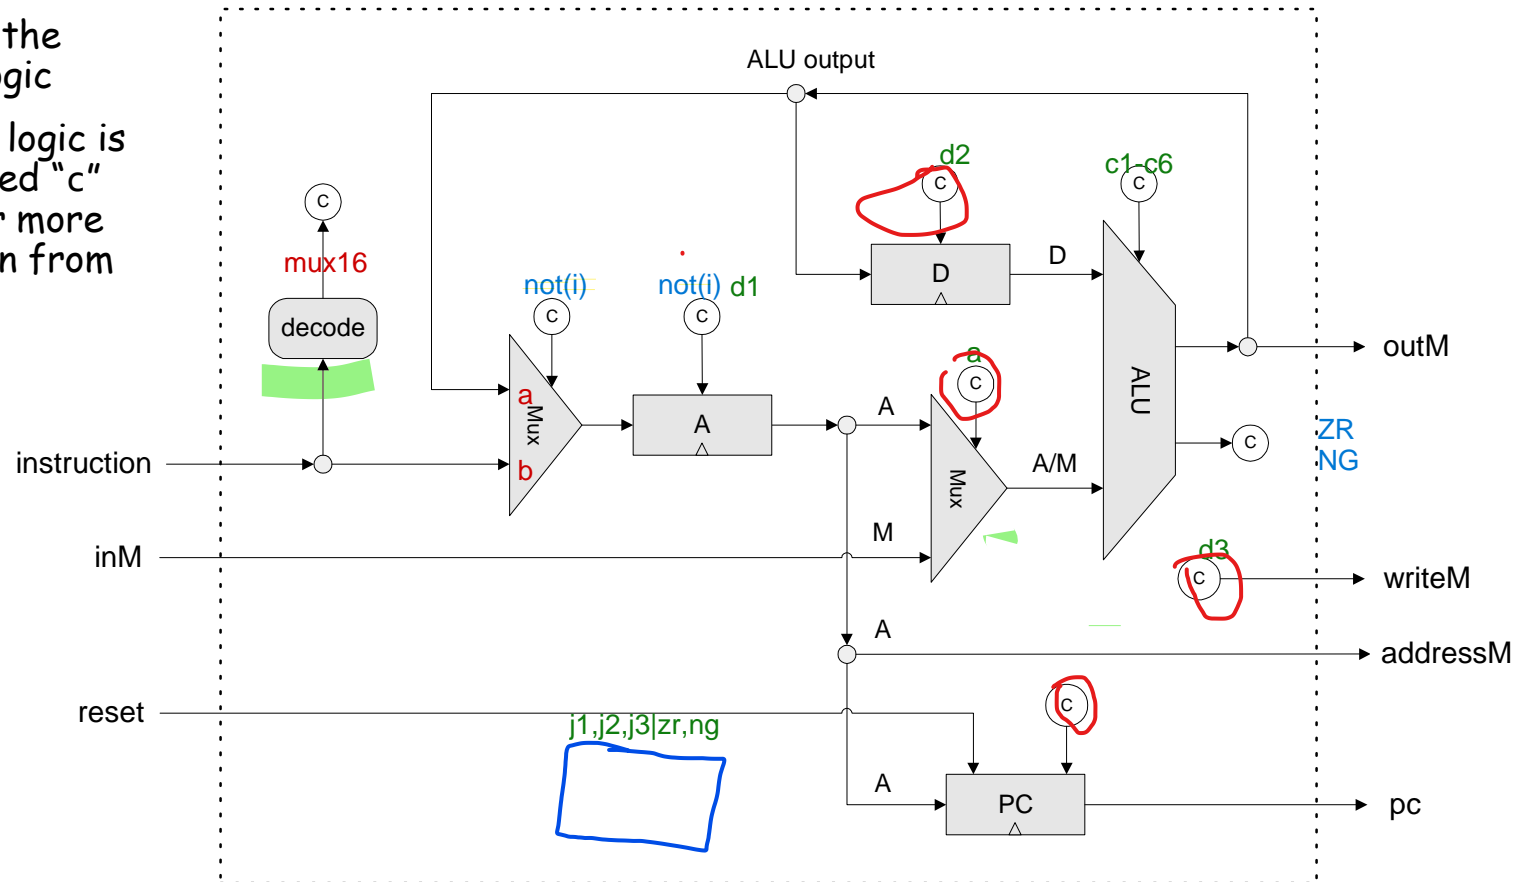
j2

j3

3, 4, 5

Chip diagram:

- Includes most of the CPU's execution logic
- The CPU's control logic is hinted: each circled "c" represents one or more control bits, taken from the instruction
- The "decode" bar does not represent a chip, but rather indicates that the instruction bits are decoded somehow.



Cycle:

- Execute
- Fetch

Execute logic:

- Decode
- Execute

Fetch logic:

If there should be a jump,
set PC to A
else set PC to PC+1


Resetting the computer:

Set reset to 1,
then set it to 0.

Lecture / construction plan

- ✓ ■ Instruction memory
- ✓ ■ Memory:
 - Data memory
 - Screen
 - Keyboard
- ✓ ■ CPU
- ➡ ■ Computer

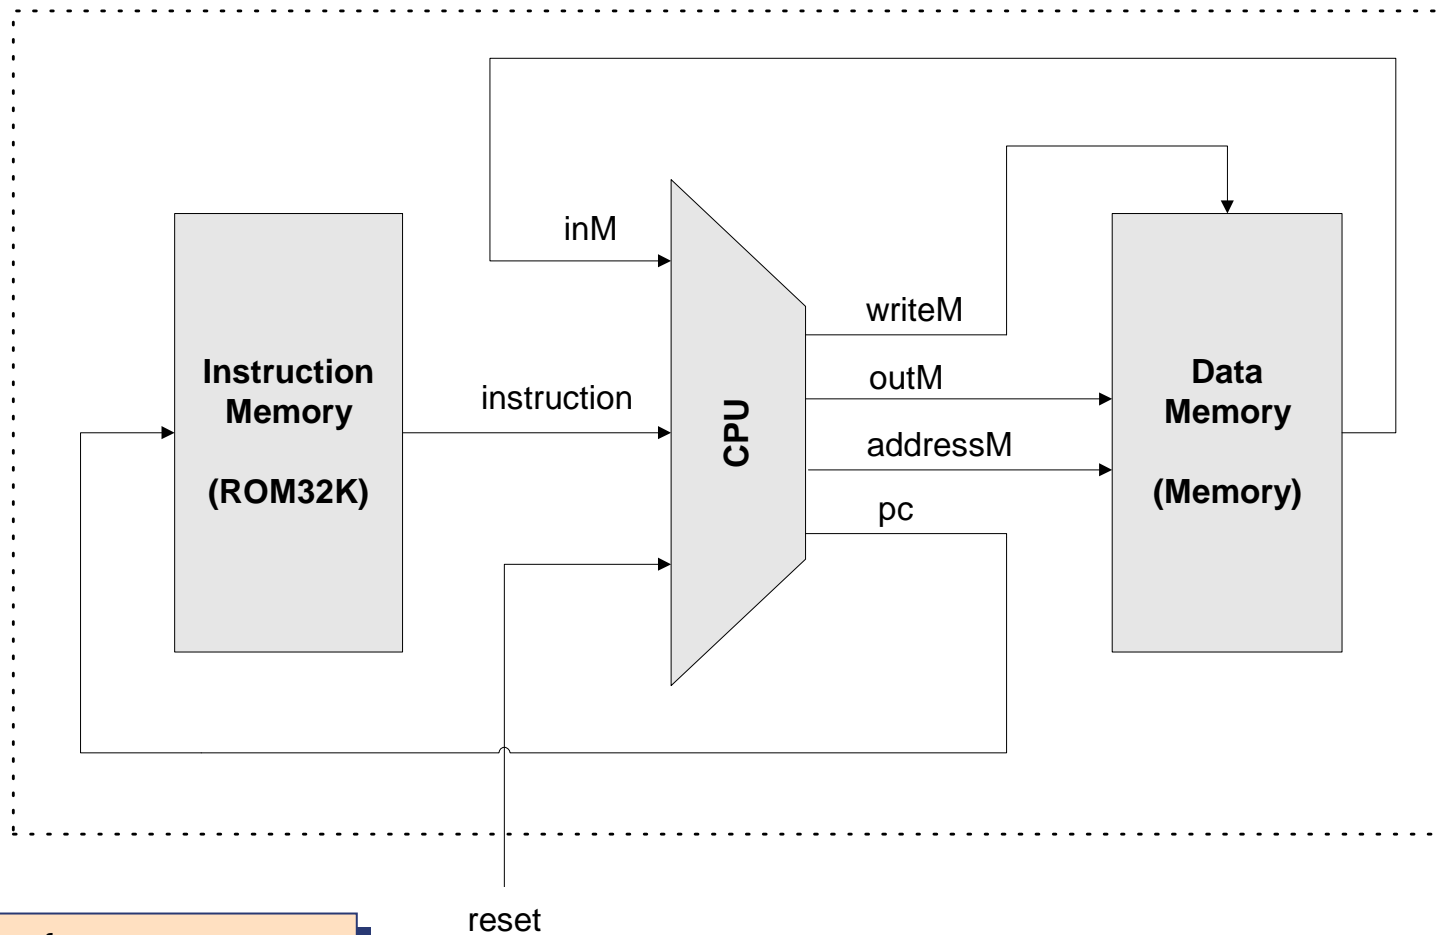
Computer-on-a-chip interface



Chip Name: Computer // Topmost chip in the Hack platform
Input: reset
Function: When reset is 0, the program stored in the computer's ROM executes. When reset is 1, the execution of the program restarts. Thus, to start a program's execution, reset must be pushed "up" (1) and "down" (0).

From this point onward the user is at the mercy of the software. In particular, depending on the program's code, the screen may show some output and the user may be able to interact with the computer via the keyboard.

Computer-on-a-chip implementation



```
CHIP Computer {  
  IN reset;  
  PARTS:  
    // implementation missing  
}
```


Implementation:

Simple, the chip-parts do all the hard work.

The spirit of things

We ascribe beauty to that which is
simple; which has no superfluous parts;
which exactly answers its end;
which stands related to all things;
which is the mean of many extremes.


(Ralph Waldo Emerson,
1803-1882)



Lecture plan



- ✓ ■ Instruction memory
- ✓ ■ Memory:
 - Data memory
 - Screen
 - Keyboard
- ✓ ■ CPU
- ✓ ■ Computer



Perspective: from here to a “real” computer

- Caching
- More I/O units
- Special-purpose processors (I/O, graphics, communications, ...)
- Multi-core / parallelism
- Efficiency
- Energy consumption considerations
- And more ...

Perspective: some issues we haven't discussed (among many)

- CISC / RISC (hardware / software trade-off)
- Hardware diversity: desktop, laptop, hand-held, game machines, ...
- General-purpose vs. embedded computers
- Silicon compilers
- And more ...